


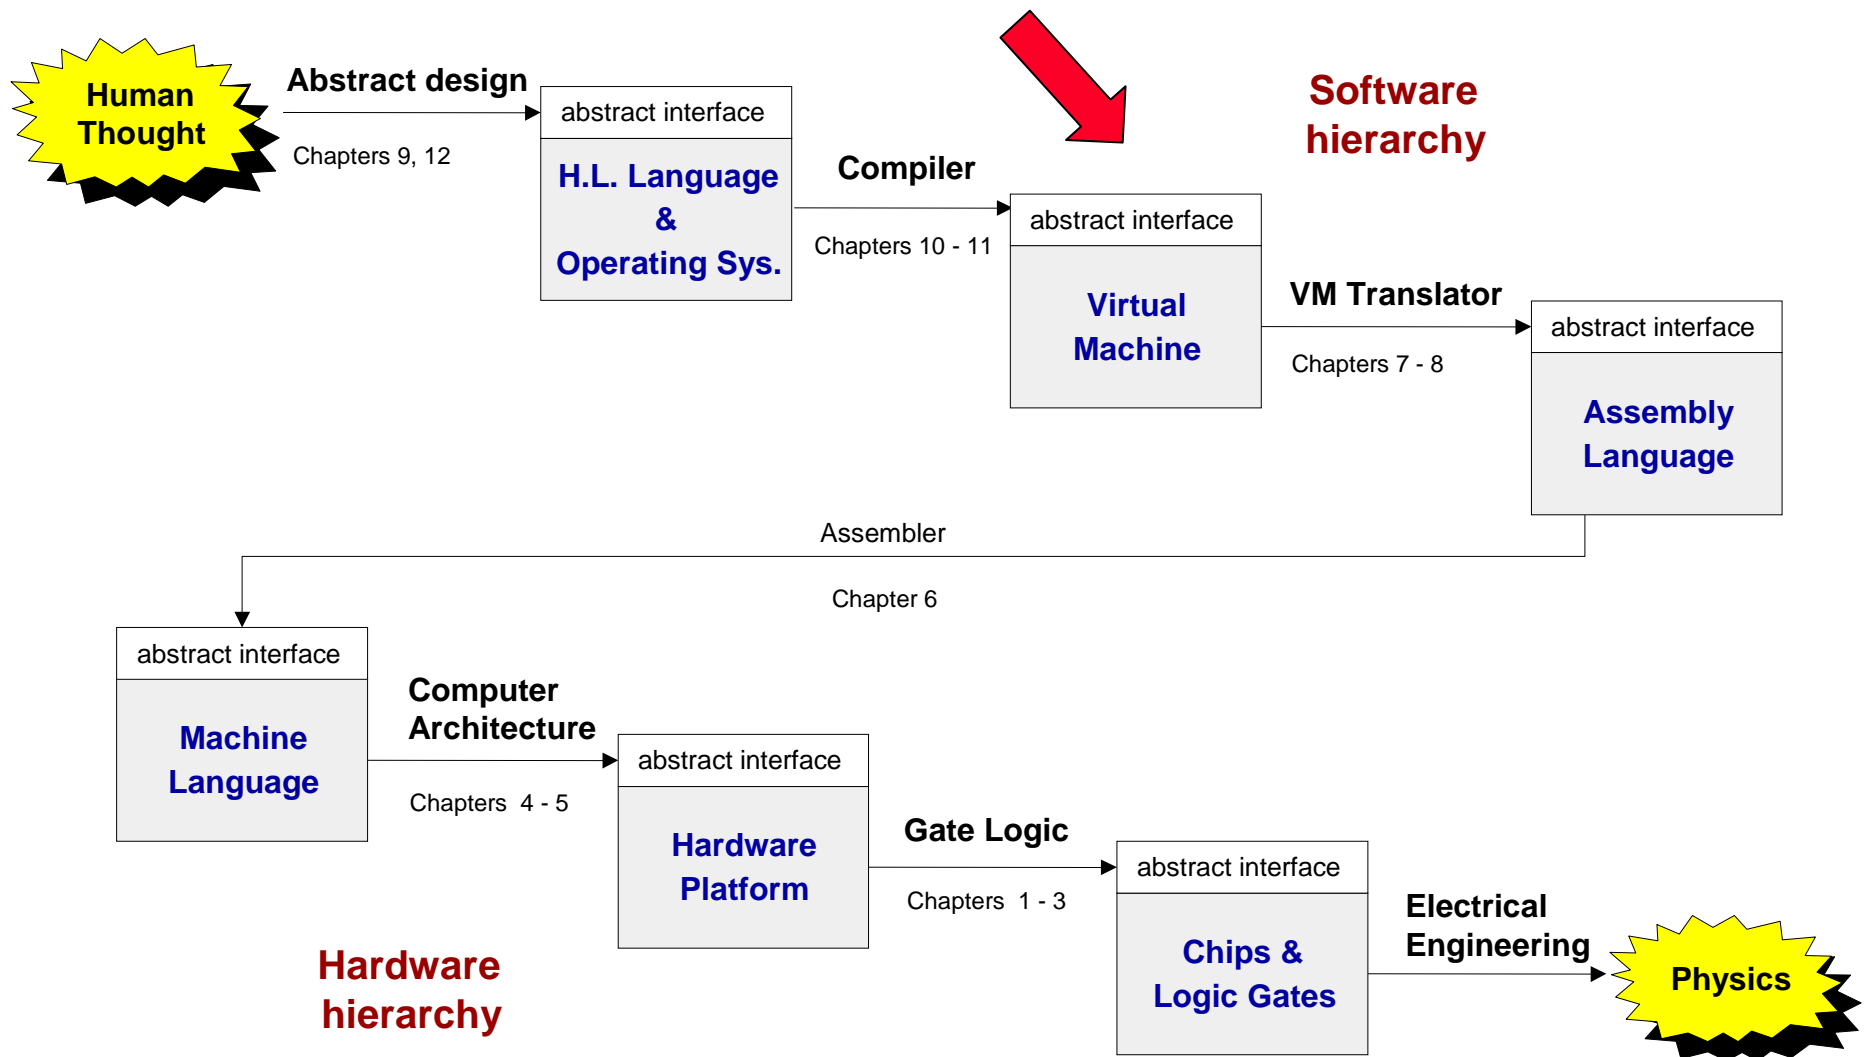
Virtual Machine I: Stack Arithmetic



Building a Modern Computer From First Principles

www.nand2tetris.org

Where we are at:



Motivation

```
class Main {
  static int x;

  function void main() {
    // Input and multiply 2 numbers
    var int a, b, x;
    let a = Keyboard.readInt("Enter a number");
    let b = Keyboard.readInt("Enter a number");
    let x = mult(a,b);
    return;
  }

  // Multiplies two numbers.
  function int mult(int x, int y) {
    var int result, j;
    let result = 0; let j = y;
    while not(j = 0) {
      let result = result + x;
      let j = j - 1;
    }
    return result;
  }
}
```

Ultimate goal:

Translate high-level programs into executable code.




Compiler

```
...
@a
M=D
@b
M=0
(LOOP)
@a
D=M
@b
D=D-A
@END
D;JGT
@j
D=M
@temp
M=D+M
@j
M=M+1
@LOOP
0;JMP
(END)
@END
0;JMP
...
```

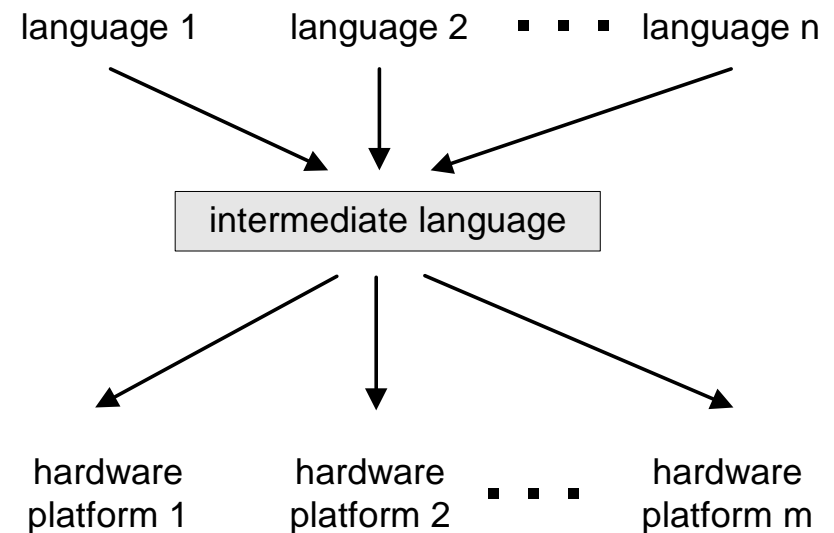
Compilation models

direct compilation:



requires $n \cdot m$ translators

2-tier compilation:

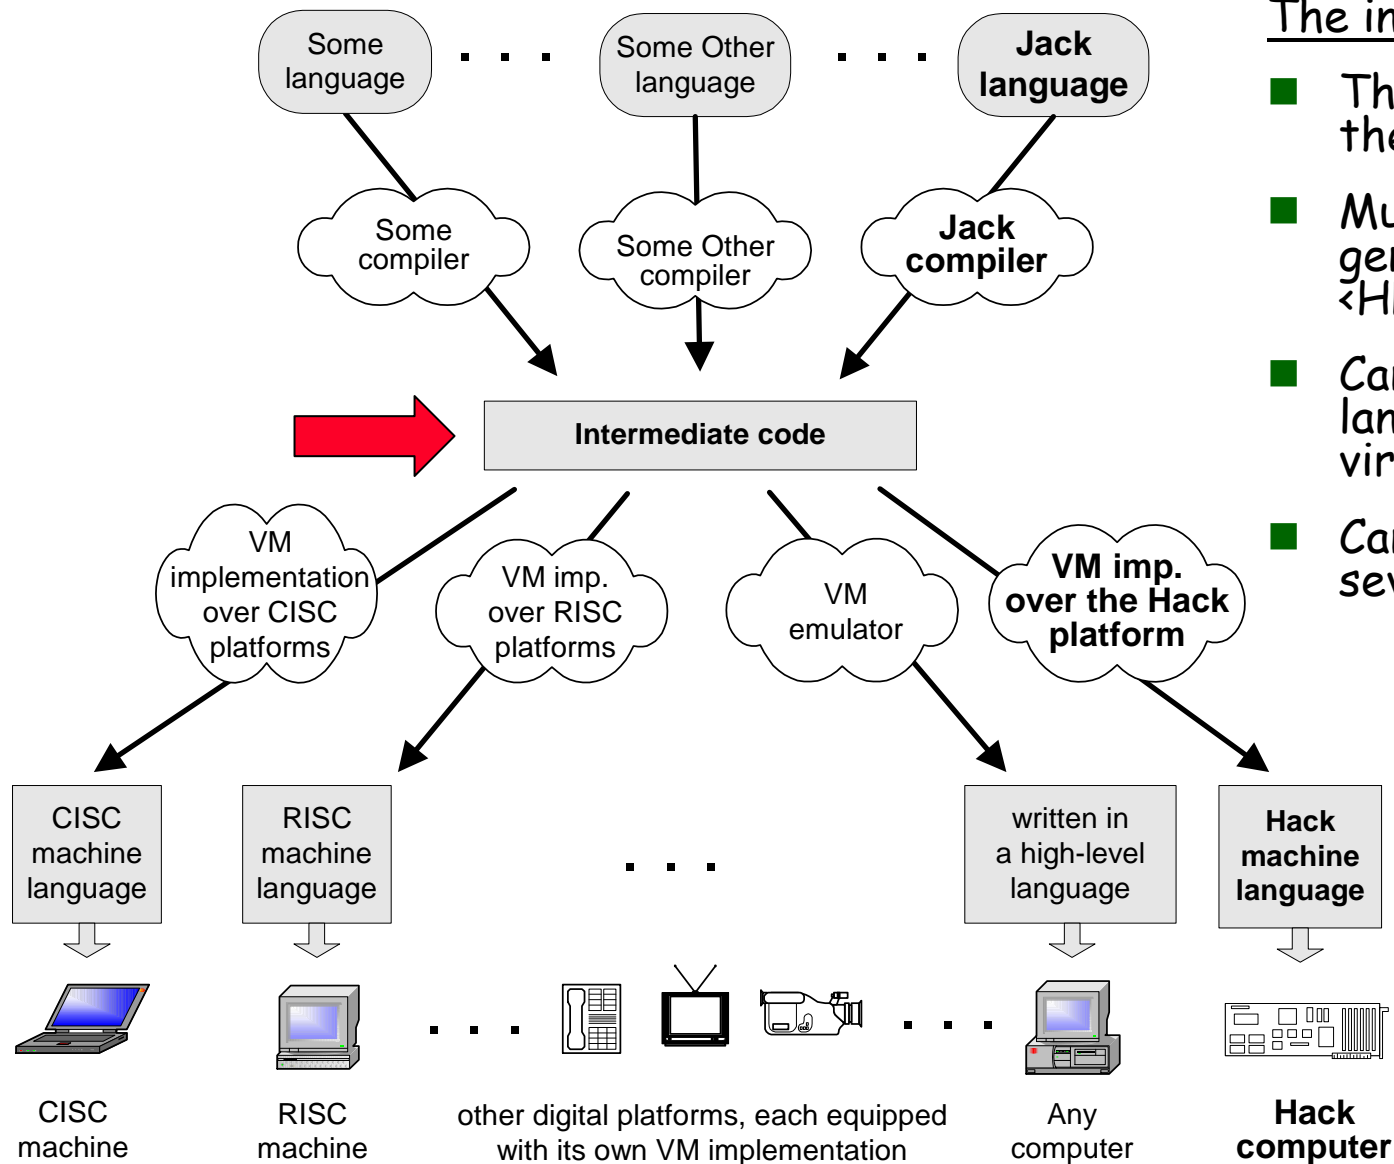


requires $n + m$ translators

Two-tier compilation:

- 1st compilation stage depends only on the details of the source language
- 2nd compilation stage depends only on the details of the target platform.

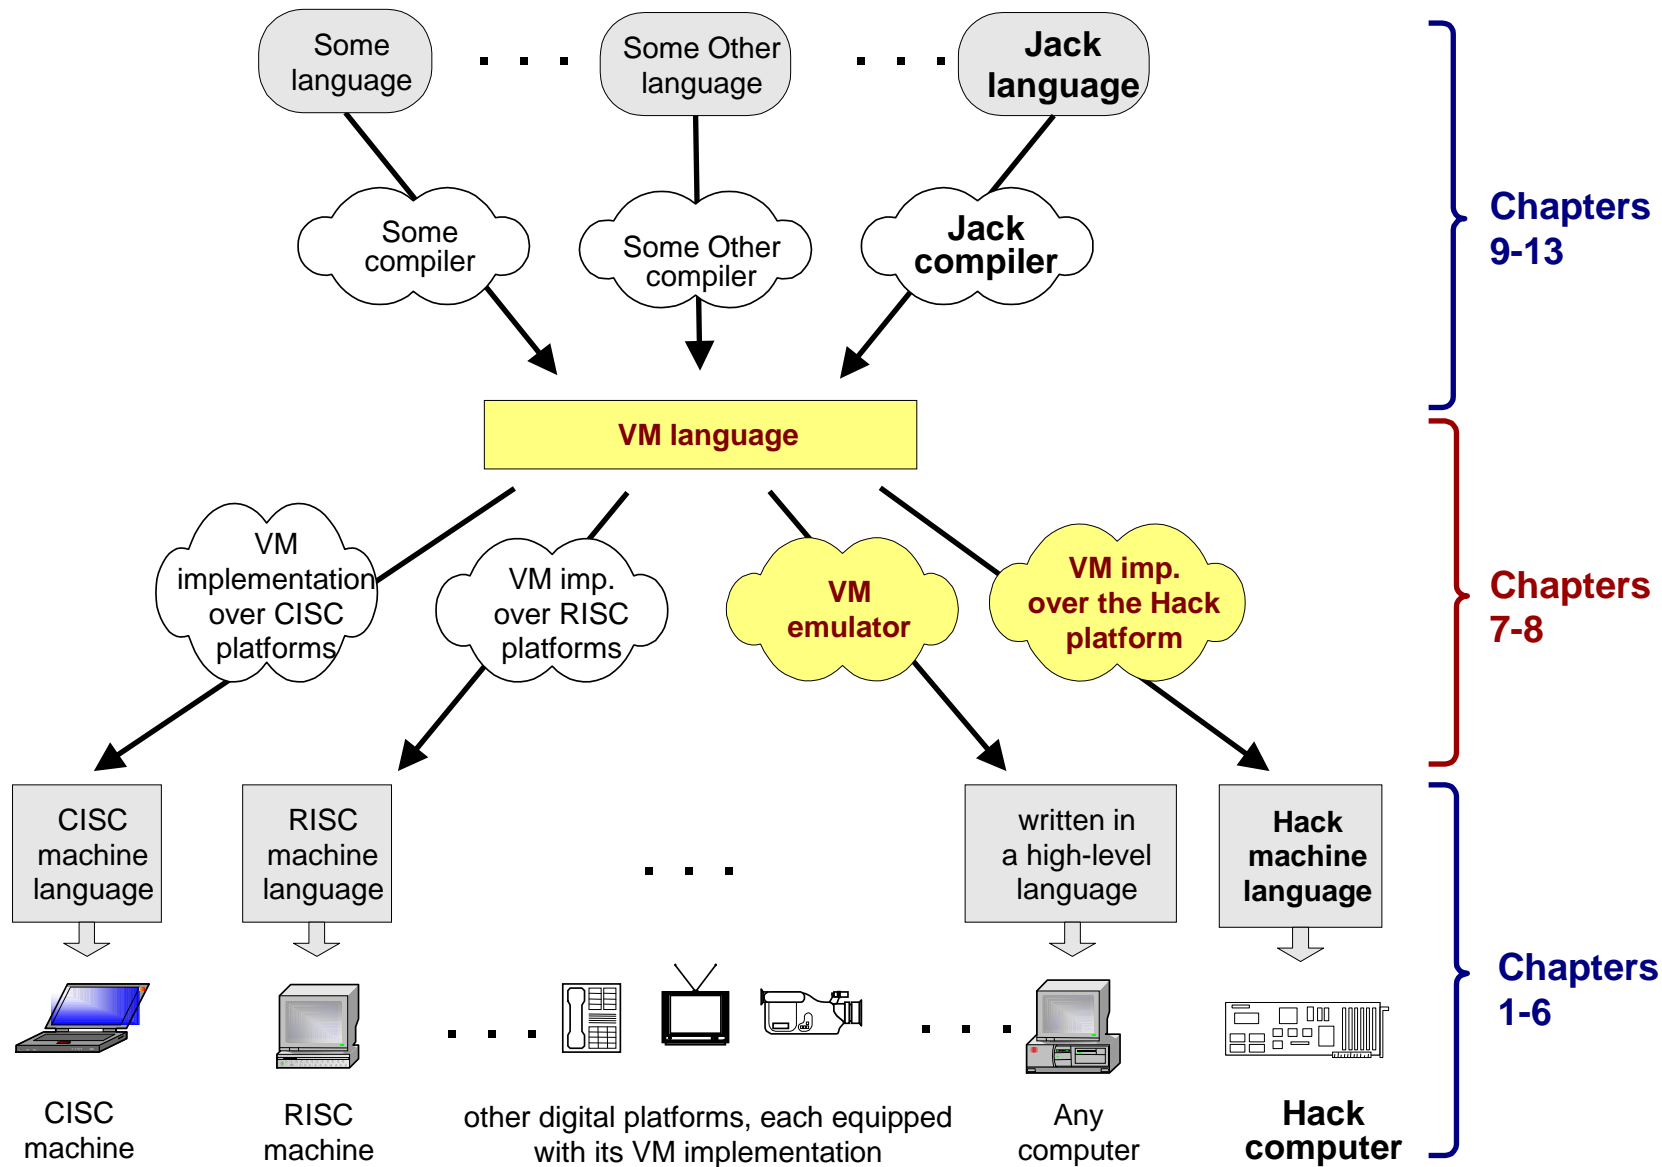
The big picture



The intermediate code:

- The interface between the 2 compilation stages
- Must be sufficiently general to support many $\langle \text{HLL}, \text{ML} \rangle$ pairs
- Can be modeled as the language of an abstract virtual machine (VM)
- Can be implemented in several different ways.

Focus of this lecture:



Lecture plan

Goal: Specify and implement a VM model and language:

Arithmetic / Boolean commands

add
sub
neg
eq
gt
lt
and
or
not

This lecture

Memory access commands

pop x (pop into x, which is a variable)
push y (y being a variable or a constant)

Program flow commands

label (declaration)
goto (label)
if-goto (label)

Next lecture

Function calling commands

function (declaration)
call (a function)
return (from a function)

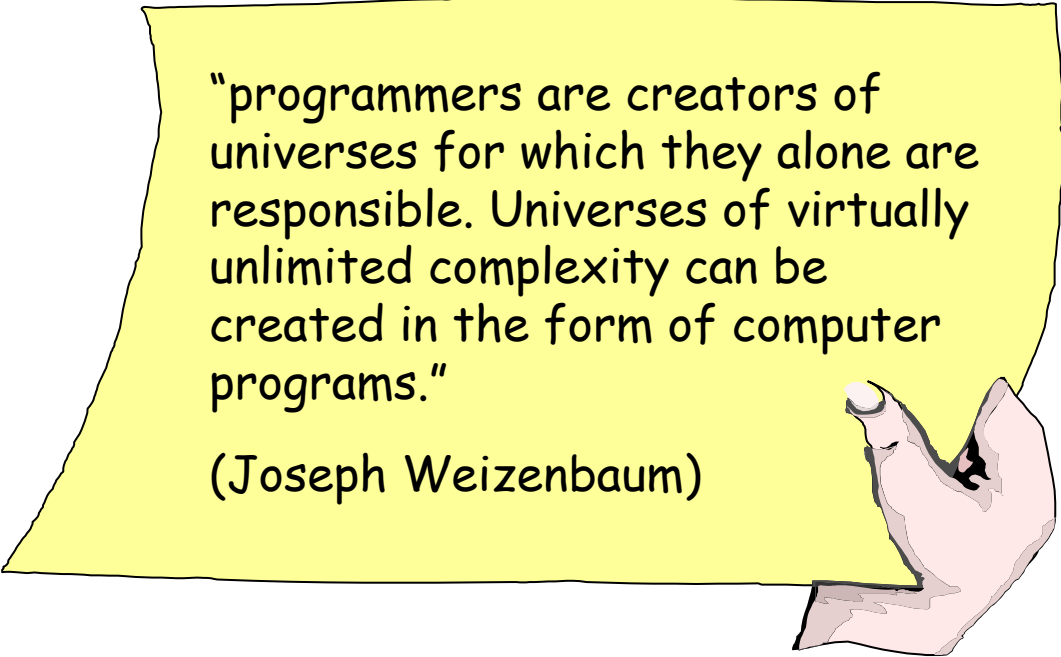
Method: (a) specify the abstraction (model's constructs and commands)
(b) propose how to implement it over the Hack platform.

The VM model and language

Perspective:

From here till the end of this and the next lecture we describe the VM model used in the Hack-Jack platform

Other VM models (like JVM/JRE and IL/CLR) are similar in spirit but different in scope and details.



"programmers are creators of universes for which they alone are responsible. Universes of virtually unlimited complexity can be created in the form of computer programs."


(Joseph Weizenbaum)

The VM model + language are an example of one such universe

View I: the VM is a machine that makes sense in its own right

View II: the VM is a convenient "in-between" representation of a computer program: an interim "station" between the high-level code and the machine level code.

Our VM is a stack-oriented machine



Our VM features a single 16-bit data type that can be used as:

- Integer
- Boolean
- Pointer

Typical operation:

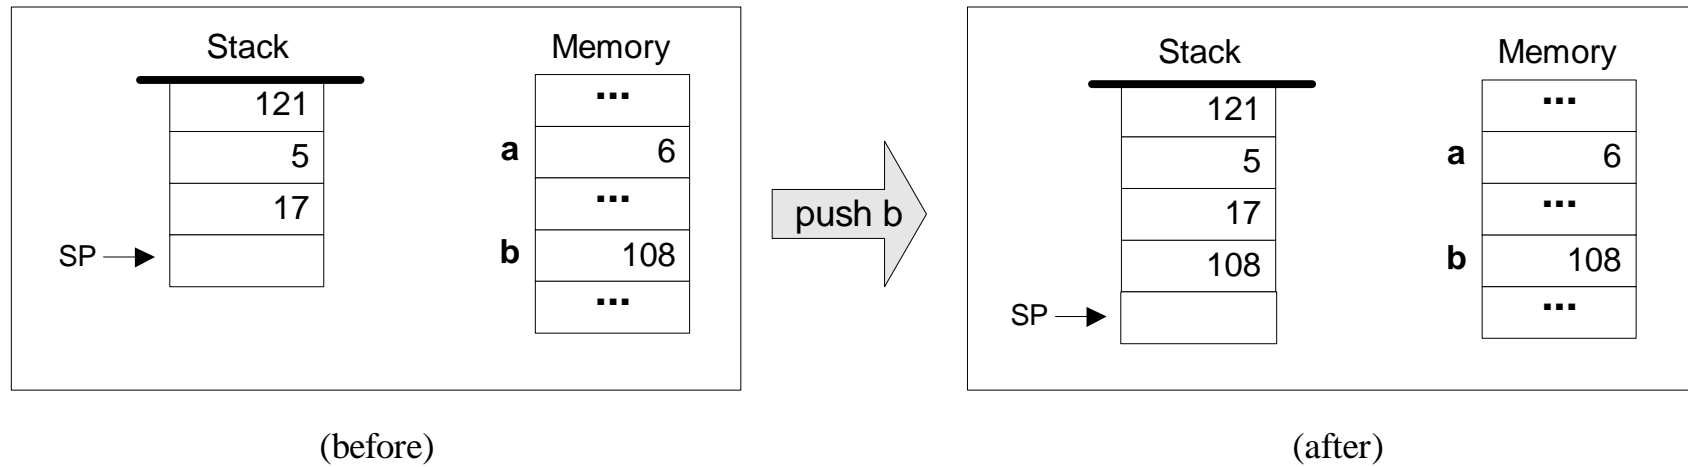
- Pops topmost values x, y from the stack
- Computes the value of some function $f(x, y)$
- Pushes the result onto the stack

(Unary operations are similar, using x and $f(x)$ instead)

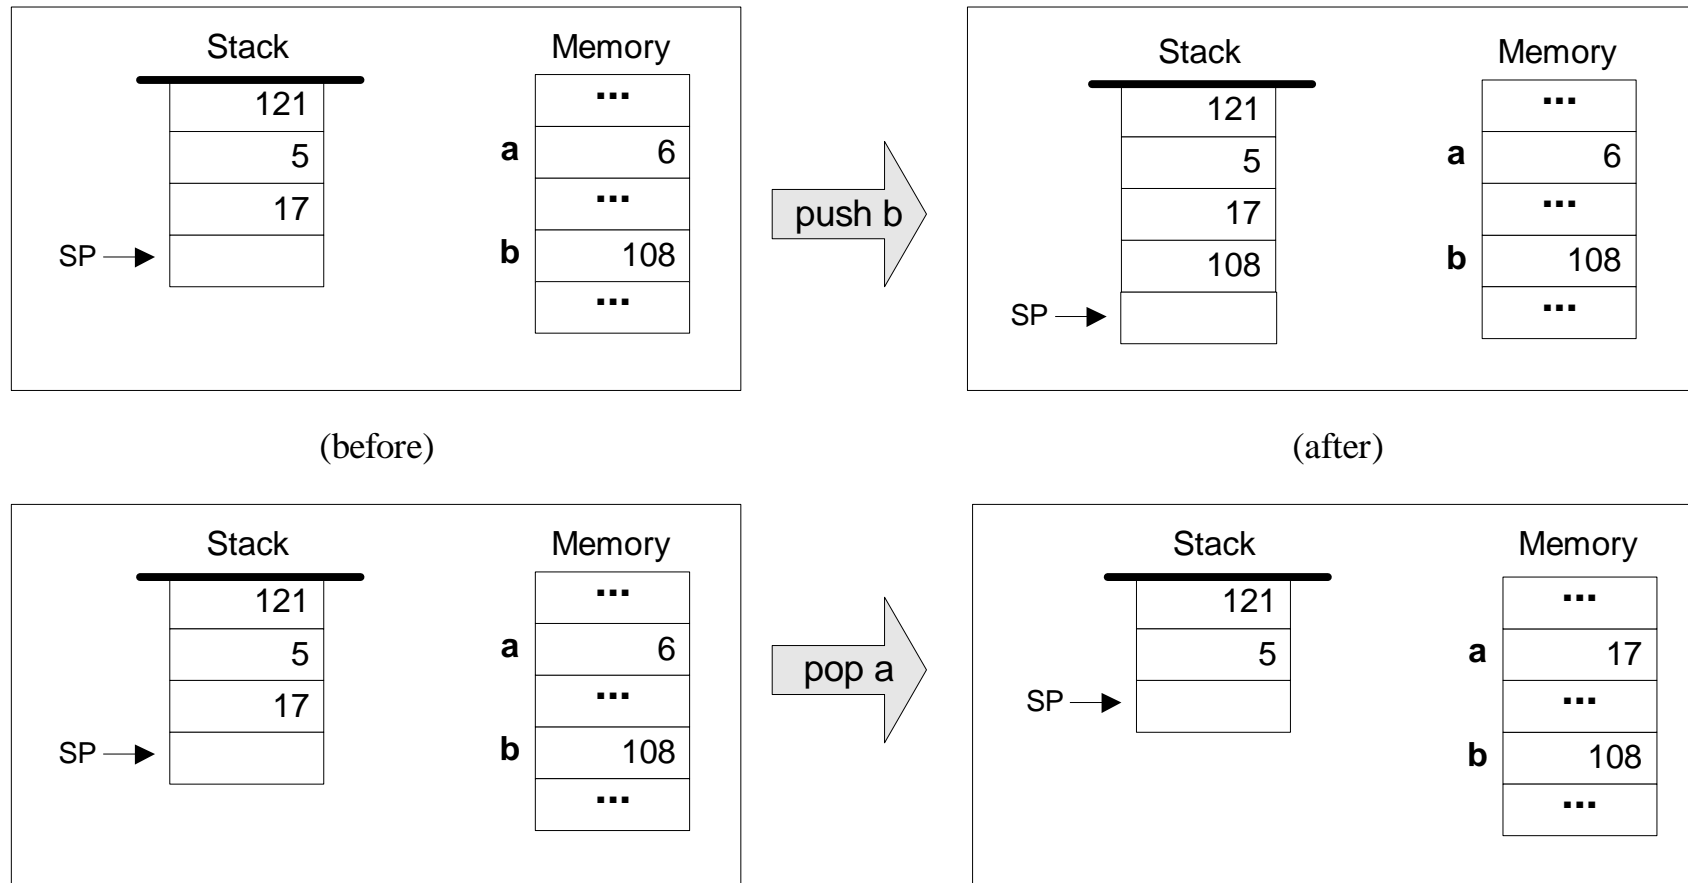
Impact: the operands are replaced with the operation's result

In general: all arithmetic and Boolean operations are implemented similarly.

Memory access (first approximation)



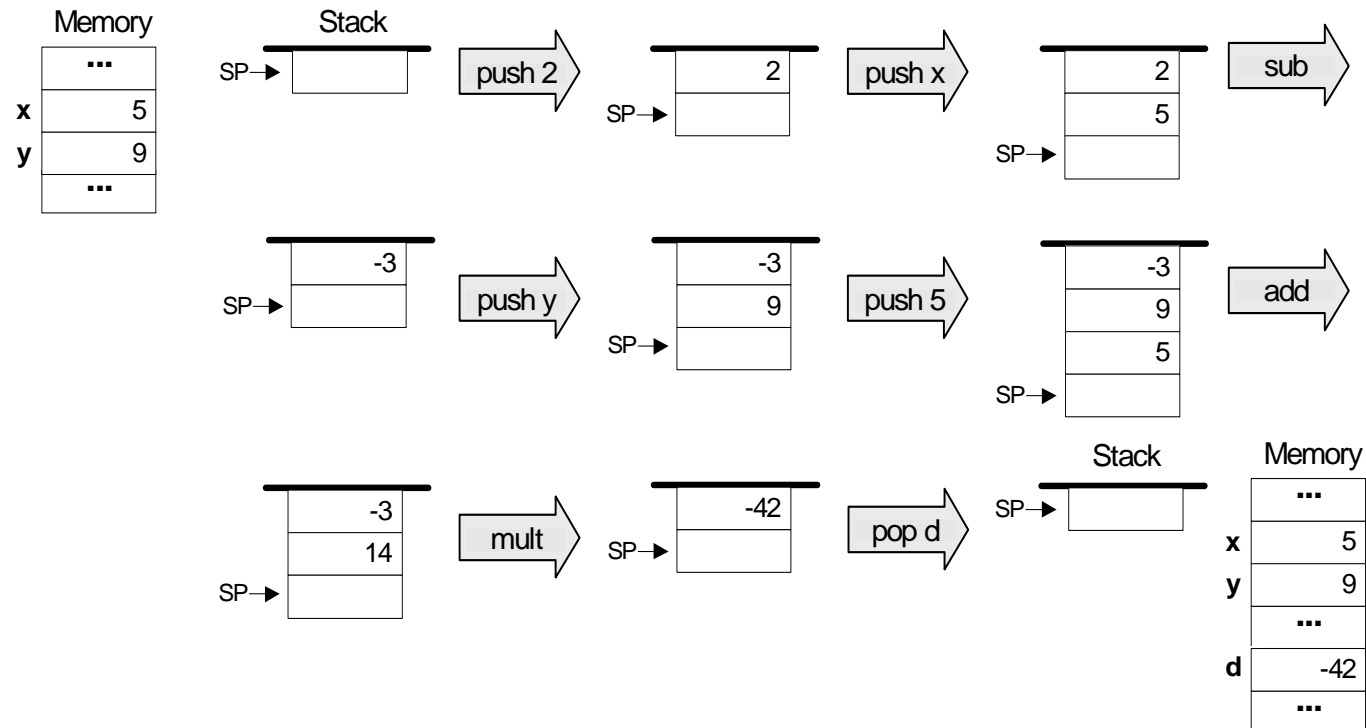
Memory access (first approximation)



- Classical data structure
- Elegant and powerful
- Several hw/sw implementation options.

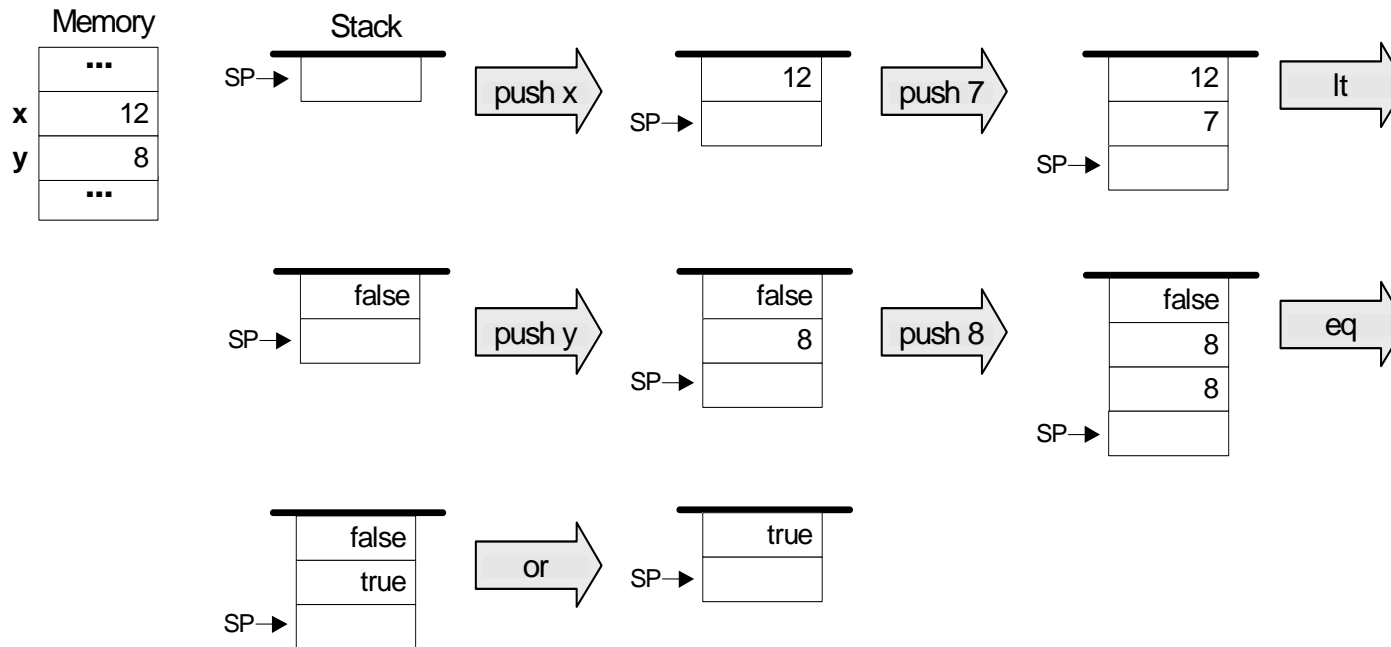
Evaluation of arithmetic expressions

```
// d=(2-x)*(y+5)
push 2
push x
sub
push y
push 5
add
mult
pop d
```




Evaluation of Boolean expressions

```
// if (x<7) or (y=8)
push x
push 7
lt
push y
push 8
eq
or
```



Arithmetic and Boolean commands (wrap-up)

Command	Return value (after popping the operand/s)	Comment
add	$x + y$	Integer addition (2's complement)
sub	$x - y$	Integer subtraction (2's complement)
neg	$-y$	Arithmetic negation (2's complement)
eq	true if $x = y$ and false otherwise	Equality
gt	true if $x > y$ and false otherwise	Greater than
lt	true if $x < y$ and false otherwise	Less than
and	$x \text{ And } y$	Bit-wise
or	$x \text{ Or } y$	Bit-wise
not	Not y	Bit-wise



Memory segments

Modern programming languages normally feature the following variable kinds:

Class level

- Static variables
- Private variables (AKA "object variables" / "fields" / "properties")

Method level:

- Local variables
- Argument variables

The VM abstraction must support (at least) these variable kinds.

In our VM model, these variables are stored in *virtual memory segments*

Specifically, our VM model consists of 8 memory segments:

`static, this, local, argument`

As well as:

`that, constant, pointer, and temp.`

Memory segments and memory access commands

So we have 8 virtual memory segments:

`static, this, local, argument, that, constant, pointer, temp`

But, at the VM level, there is no need to differentiate among the different roles of these segments

As far as VM programming commands go, a segment is a segment is a segment;
to access a particular segment entry, use the syntax:

Memory access command format:

`pop segment i`

`push segment i`

- These commands are used instead of `pop x` and `push y`, as shown in previous slides, which was a conceptual simplification
- The meaning of the different eight segments will become important when we'll talk about the compiler.

VM programming


- VM programs are normally written by *compilers*, not by humans
- But, for a human to write or optimize a compiler, the human must first understand the spirit of VM programming
- So, here is an example.

Arithmetic example


High-level code

```
function mult (x,y) {  
  int result, j;  
  result = 0;  
  j = y;  
  while ~(j == 0) {  
    result = result + x;  
    j = j - 1;  
  }  
  return result;  
}
```

Just after mult(7,3) is entered:



Just after mult(7,3) returns:



VM code (first approx.)

```
function mult(x,y)  
  push 0  
  pop result  
  push y  
  pop j  
label loop  
  push j  
  push 0  
  eq  
  if-goto end  
  push result  
  push x  
  add  
  pop result  
  push j  
  push 1  
  sub  
  pop j  
  goto loop  
label end  
  push result  
  return
```

VM code

```
function mult 2  
  push constant 0  
  pop local 0  
  push argument 1  
  pop local 1  
label loop  
  push local 1  
  push constant 0  
  eq  
  if-goto end  
  push local 0  
  push argument 0  
  add  
  pop local 0  
  push local 1  
  push constant 1  
  sub  
  pop local 1  
  goto loop  
label end  
  push local 0  
  return
```

Lecture plan

Goal: Specify and implement a VM model and language

Arithmetic / Boolean commands

`add`

`sub`

`neg`

`eq`

`gt`

`lt`

`and`

`or`

`not`

Memory access commands

`pop segment i`

`push segment i`

This lecture

Program flow commands

`label` (declaration)

`goto` (label)

`if-goto` (label)

Next lecture

Function calling commands

`function` (declaration)

`call` (a function)

`return` (from a function)

Method: (a) specify the abstraction (model's constructs and commands)
→ (b) propose how to implement it over the Hack platform.

Implementation

VM implementation options:

- Software-based (emulate the VM using, say, Java)
- Translator-based (translate VM programs into, say, the Hack language)
- Hardware-based (realize the VM using dedicated memory and registers)

Well-known translator-based implementations:

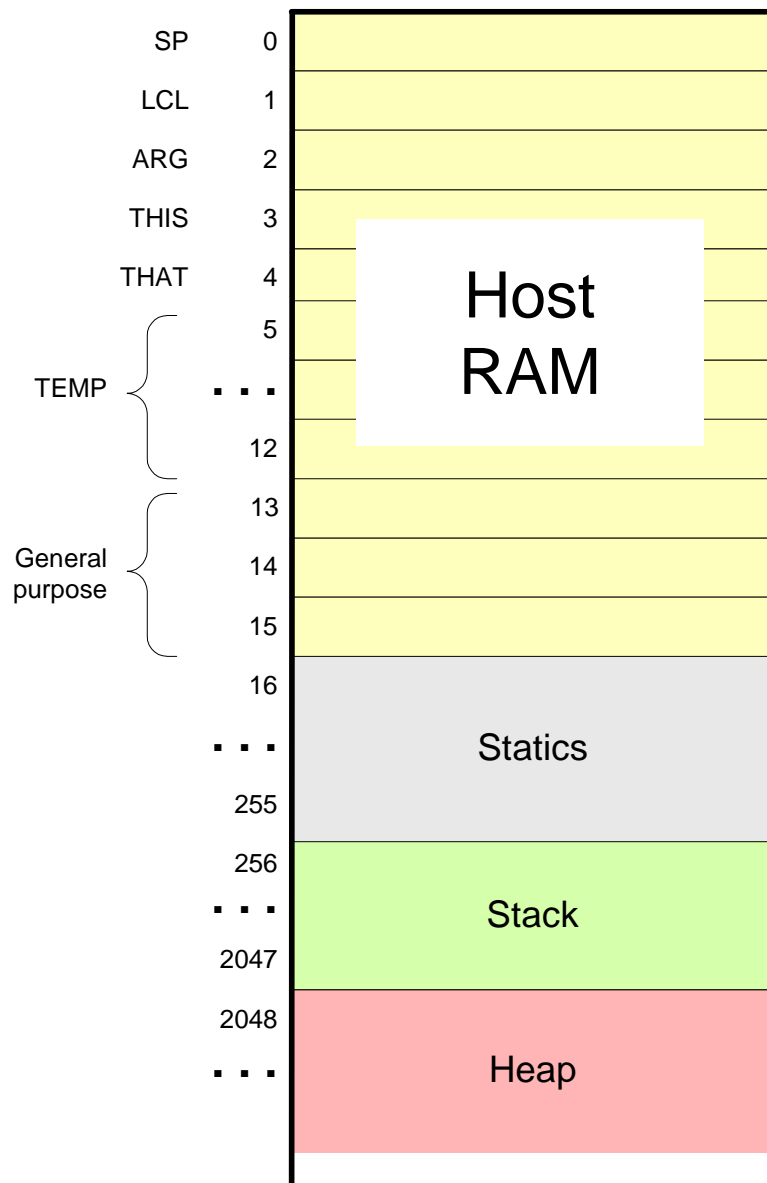
- JVM (runs bytecode programs in the Java platform)
- CLR (runs IL programs in the .NET platform).

Our VM emulator (part of the course software suite)

The screenshot shows the Virtual Machine Emulator (1.4b3) interface. The title bar indicates the file path is G:\examples\add. The menu bar includes File, View, Run, and Help. The toolbar contains icons for file operations and execution controls (Slow, Fast, Animate, View, Format). The main window is divided into several panels:

- Program:** A list of instructions. Instruction 11, 'add', is highlighted in yellow and labeled 'VM code'.
- Static:** A table for static variables.
- Local:** A table for local variables. It is labeled 'virtual memory segments'.
- Argument:** A table for arguments.
- This:** A table for 'this' pointer.
- That:** A table for 'that' pointer.
- Temp:** A table for temporary variables.
- Global Stack:** A table showing memory addresses and values. It is labeled 'global stack'.
- RAM:** A table showing memory addresses and values. It is labeled 'host RAM'. A note in blue text states: '(the RAM is not part of the VM)'. The SP register is highlighted in yellow.
- Call Stack:** A list of active functions: Sys.init, Main.main, and Main.add.
- Stack:** A table showing the current stack state. It is labeled 'working stack'.
- emulator controls:** A box containing the 'Repeat' loop and 'vmstep;' instruction, labeled 'default test script'.

Standard VM implementation on the Hack platform



The challenge: (i) map the VM constructs on the host RAM, and (ii) given this mapping, figure out how to implement each VM command using assembly commands that operate on the RAM

- **local, argument, this, that:** mapped on the heap. The base addresses of these segments are kept in **LCL, ARG, THIS, THAT**. Access to the i -th entry of a segment is implemented by accessing the segment's (base + i) word in the RAM

- **static:** static variable number j in a VM file f is implemented by the assembly language symbol $f.j$ (and recall that the assembler maps such symbols to the RAM starting from address 16)

- **constant:** truly a virtual segment: Access to **constant i** is implemented by supplying the constant i

- **pointer:** used to align **this** and **that** with memory blocks on the heap

Exercise: given the above game rules, write the Hack commands that implement, say, **push constant 5** and **pop local 2**.

Proposed VM translator implementation: Parser module

Parser: Handles the parsing of a single .vm file, and encapsulates access to the input code. It reads VM commands, parses them, and provides convenient access to their components. In addition, it removes all white space and comments.


Routine	Arguments	Returns	Function
Constructor	Input file / stream	--	Opens the input file/stream and gets ready to parse it.
hasMoreCommands	--	boolean	Are there more commands in the input?
advance	--	--	Reads the next command from the input and makes it the current command. Should be called only if hasMoreCommands () is true. Initially there is no current command.
commandType	--	C_ARITHMETIC , C_PUSH, C_POP , C_LABEL, C_GOTO , C_IF, C_FUNCTION , C_RETURN, C_CALL	Returns the type of the current VM command. C_ARITHMETIC is returned for all the arithmetic commands.
arg1	--	string	Returns the first argument of the current command. In the case of C_ARITHMETIC , the command itself (add, sub, etc.) is returned. Should not be called if the current command is C_RETURN.
arg2	--	int	Returns the second argument of the current command. Should be called only if the current command is C_PUSH, C_POP, C_FUNCTION, or C_CALL.

Proposed VM translator implementation: CodeWriter module

CodeWriter: Translates VM commands into Hack assembly code.			
Routine	Arguments	Returns	Function
Constructor	Output file / stream	--	Opens the output file/stream and gets ready to write into it.
setFileName	fileName (string)	--	Informs the code writer that the translation of a new VM file is started.
writeArithmetic	command (string)	--	Writes the assembly code that is the translation of the given arithmetic command.
WritePushPop	Command (C_PUSH or C_POP), segment (string), index (int)	--	Writes the assembly code that is the translation of the given command, where command is either C_PUSH or C_POP.
Close	--	--	Closes the output file.
Comment: More routines will be added to this module in chapter 8.			

Perspective

- In this lecture we began the process of building a compiler
- Modern compiler architecture:
 - Front-end (translates from high level language to a VM language)
 - Back-end (translates from the VM language to the machine language of some target hardware platform)
- Brief history of virtual machines:
 - 1970's: p-Code
 - 1990's: Java's JVM
 - 2000's: Microsoft .NET
- A full blown VM implementation typically includes a common software library (can be viewed as a mini, portable OS).
- We will build such a mini OS later in the course.



The road ahead

Tasks:

- Complete the VM specification and implementation (chapters 7,8)
- Introduce Jack, a high-level programming language (chapter 9)
- Build a compiler for it (chapters 10,11)
- Finally, build a mini-OS, i.e. a run-time library (chapter 12).

Conceptually similar to:

- JVM
- Java
- Java compiler
- JRE



And to:

- CLR
- C#
- C# compiler
- .NET base class library

